


Fault Tolerance - midterm

Nikhil, Jessica, Fabian

Motivation

- Fault Tolerance - ability to continue functioning in the event of failure.

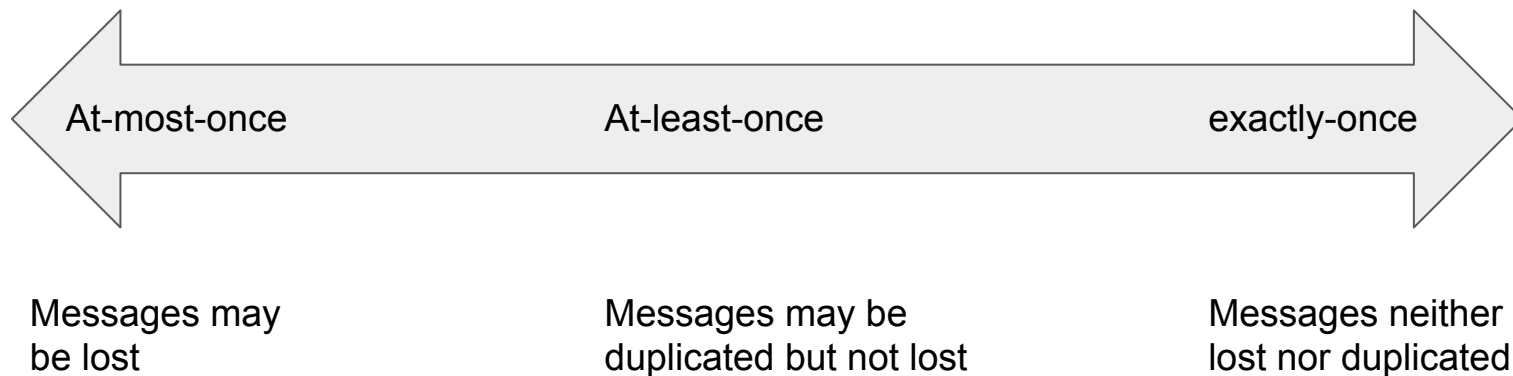
- Objective: provide dependability



- Failure masking Techniques:

- Replication
- Checkpointing


Data handling guarantees



Flink fault tolerance

How to ensure exactly-once semantics for the state?

periodically snapshot the memory



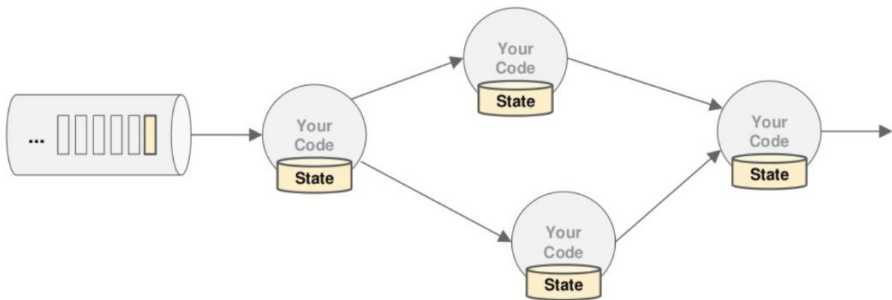
Recovery: restore snapshot and replay events since snapshot



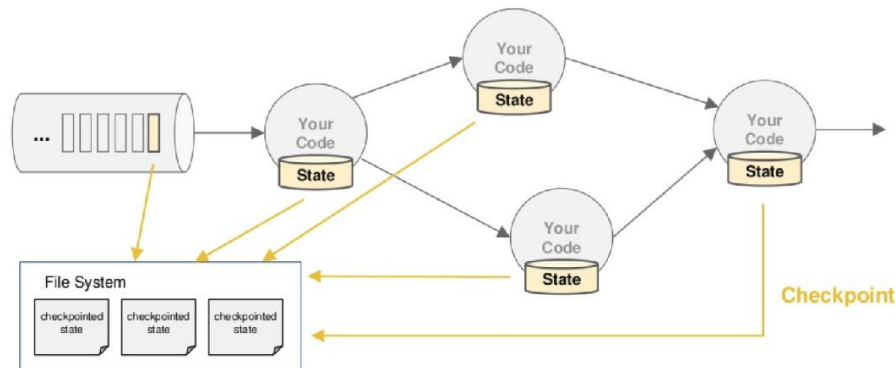
Flink checkpointing

- Draws consistent snapshots of the distributed data stream and operator state.

Consistent snapshotting:

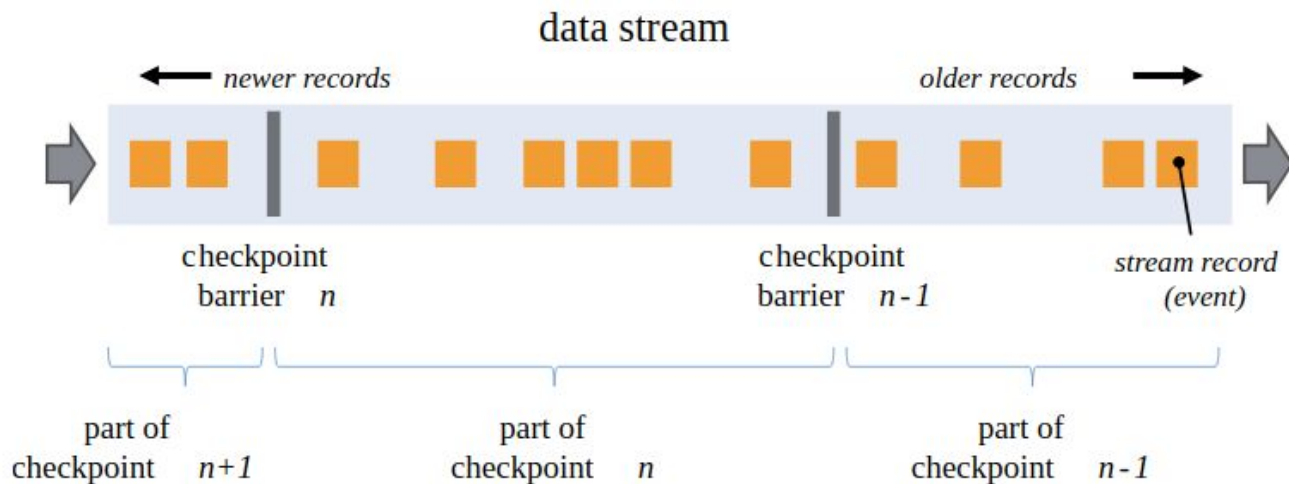


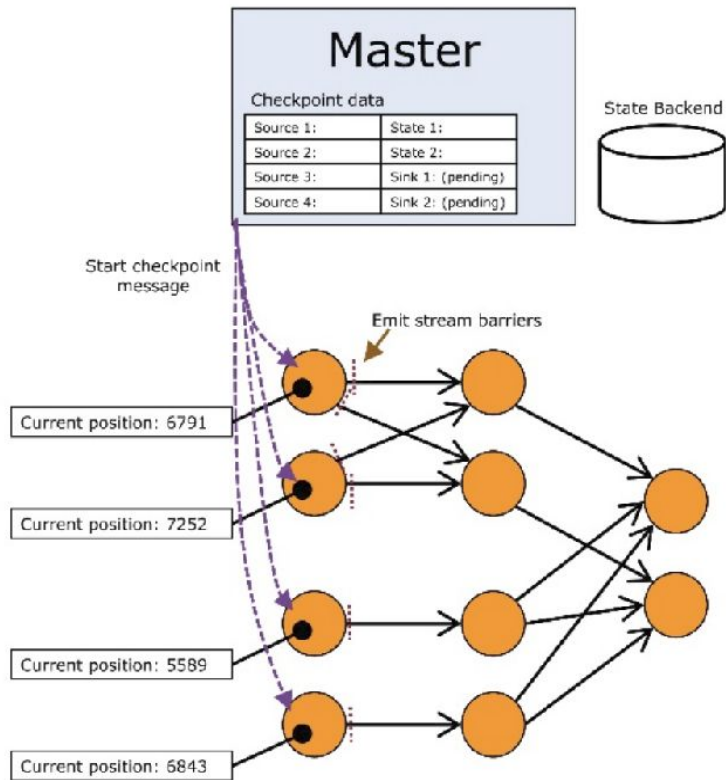
Consistent snapshotting:



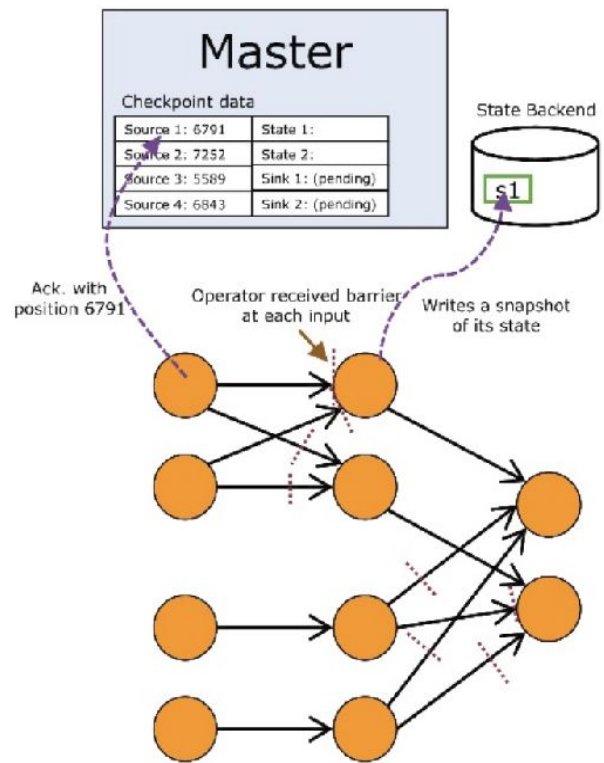
Checkpoint barriers

- Markers for checkpoints
- Injected into Datastream







Starting
Checkpoint




Checkpoint
in Progress




Ack. with pointer
to state



Checkpoint
in Progress



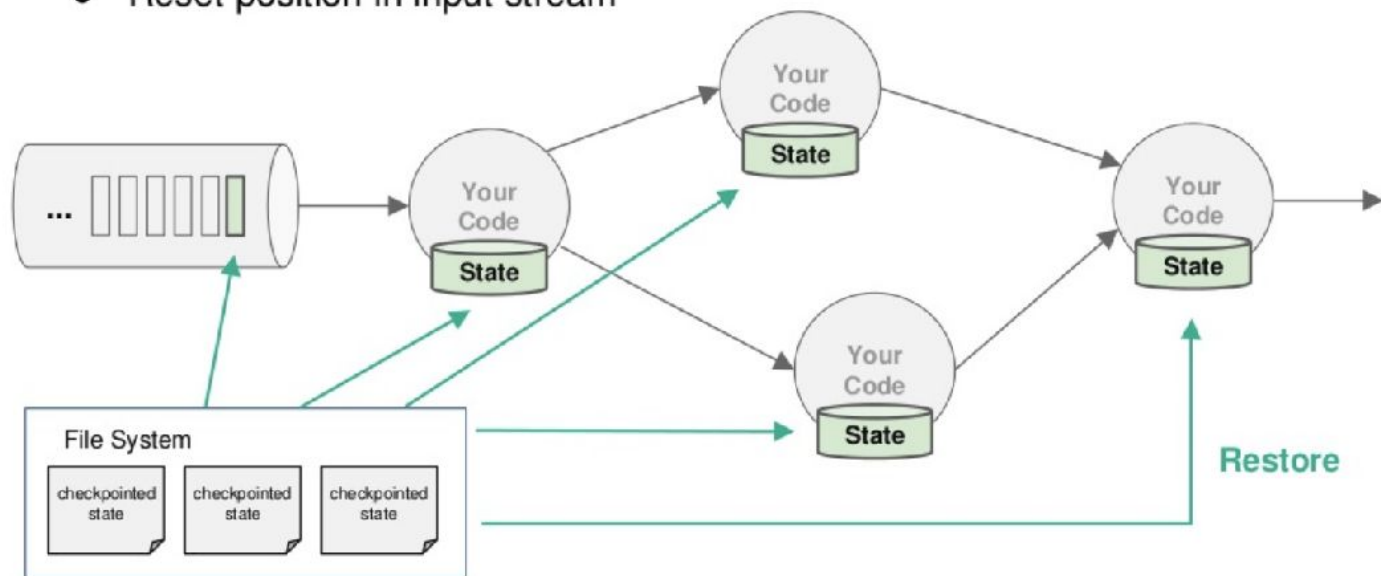
Sink acknowledges
checkpoint after
receiving all barriers




Checkpoint
Completed

Flink Recovery


- Recover all embedded state
- Reset position in input stream




Single Point of Failure



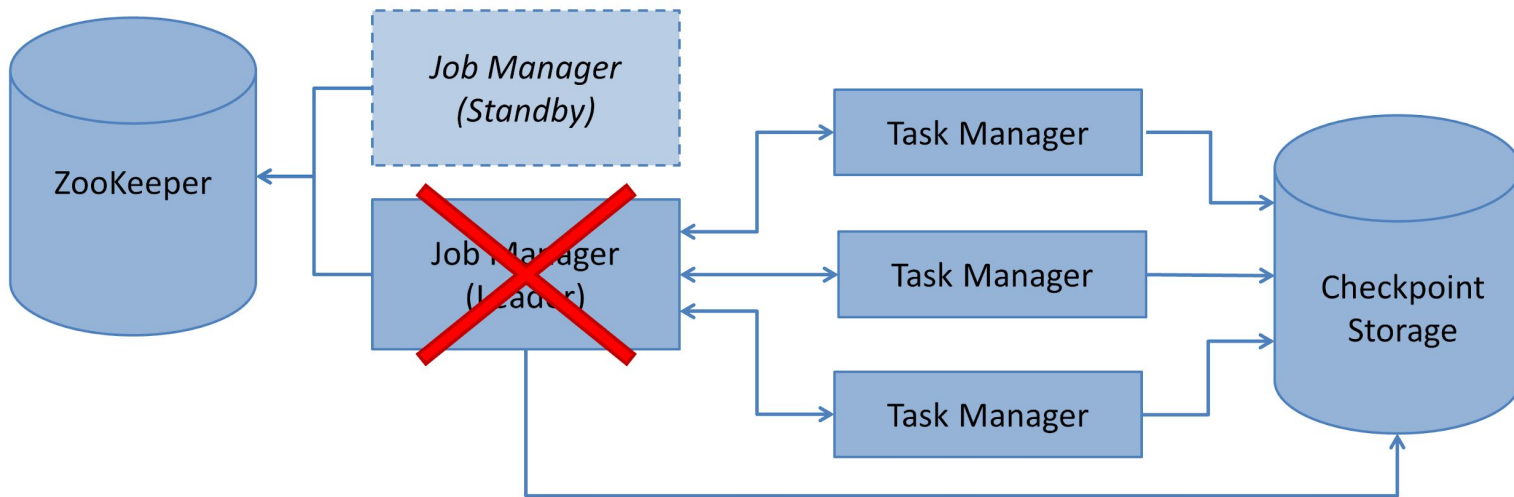
Single Point of Failure



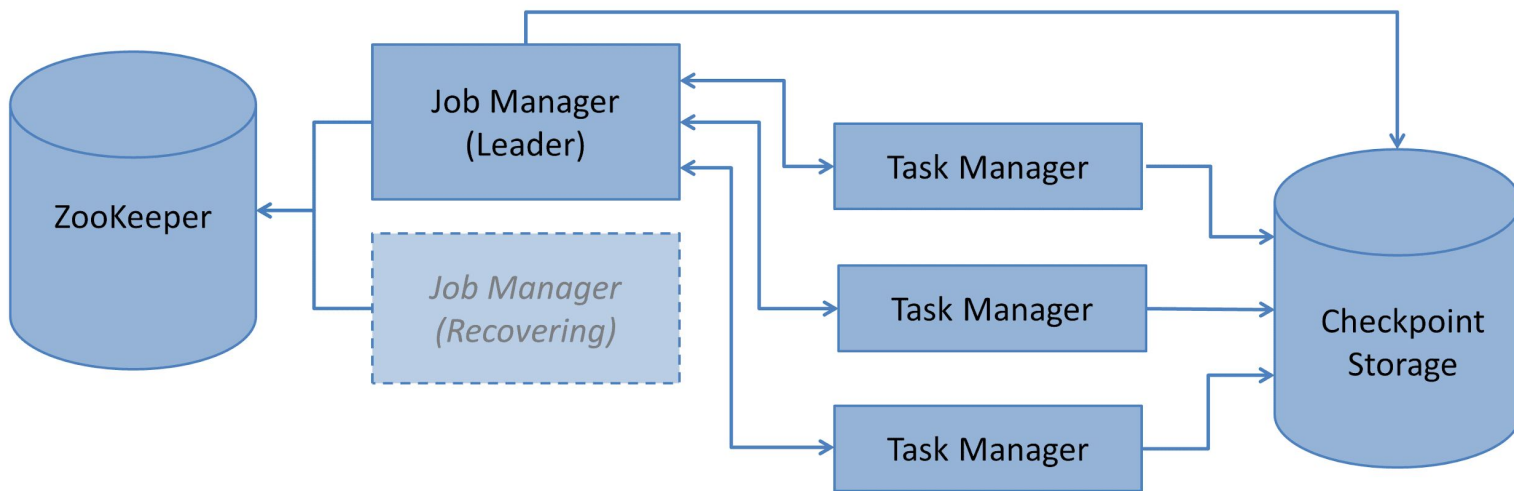
High Availability Mode (Standalone)




High Availability Mode (Standalone)




High Availability Mode (Standalone)



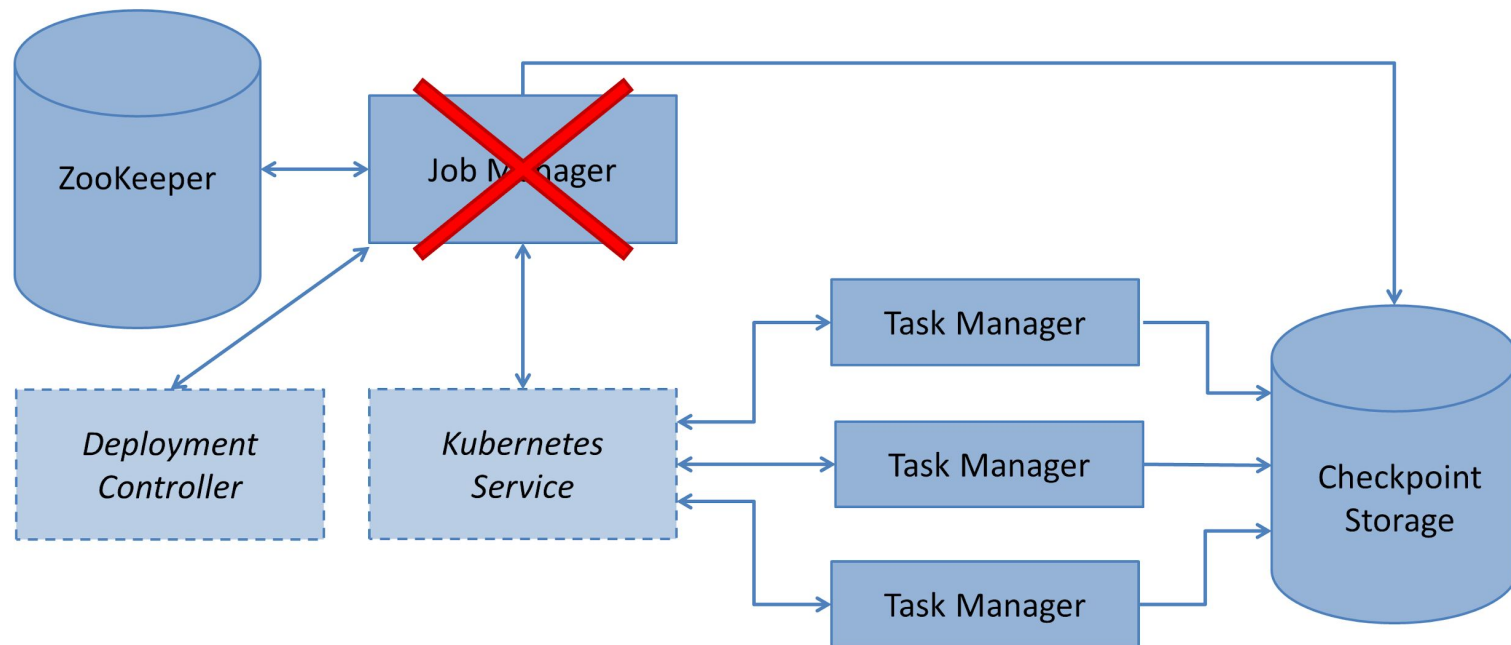
High Availability Mode (Standalone)



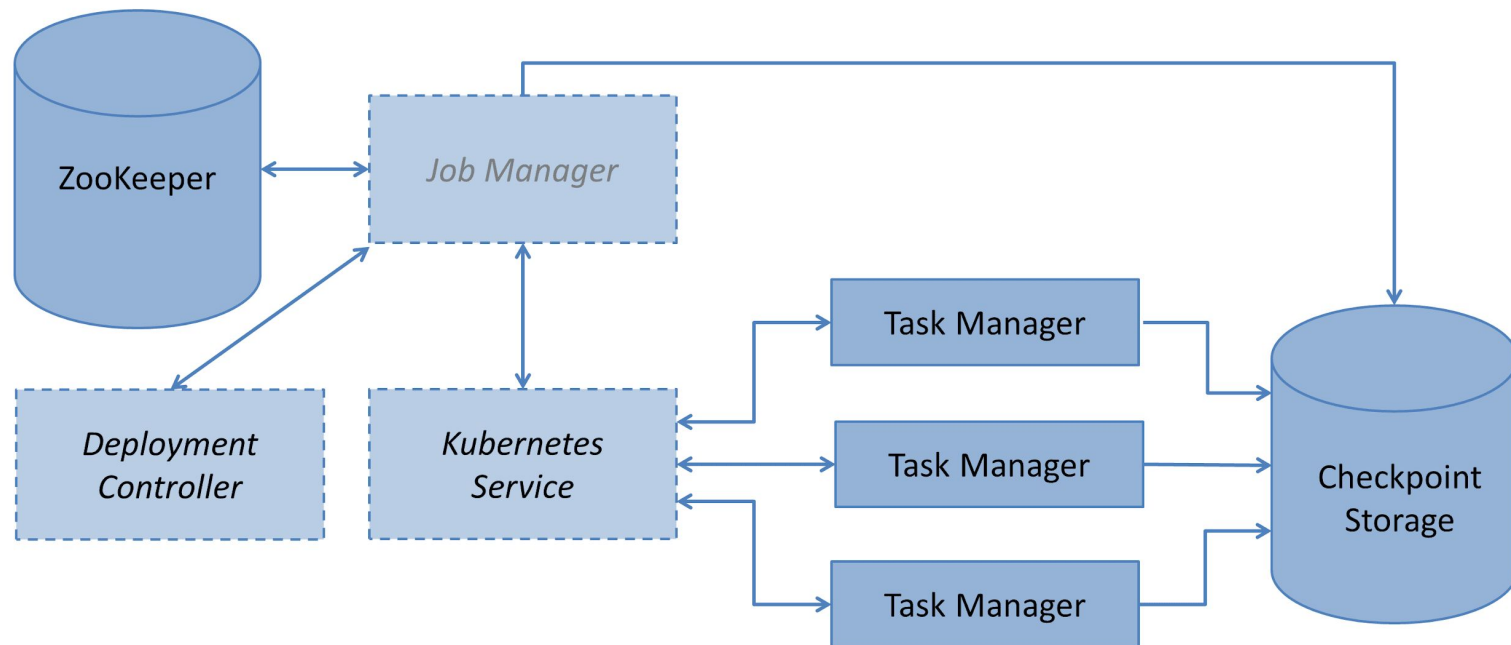
High Availability Mode (w/ Kubernetes or Yarn)



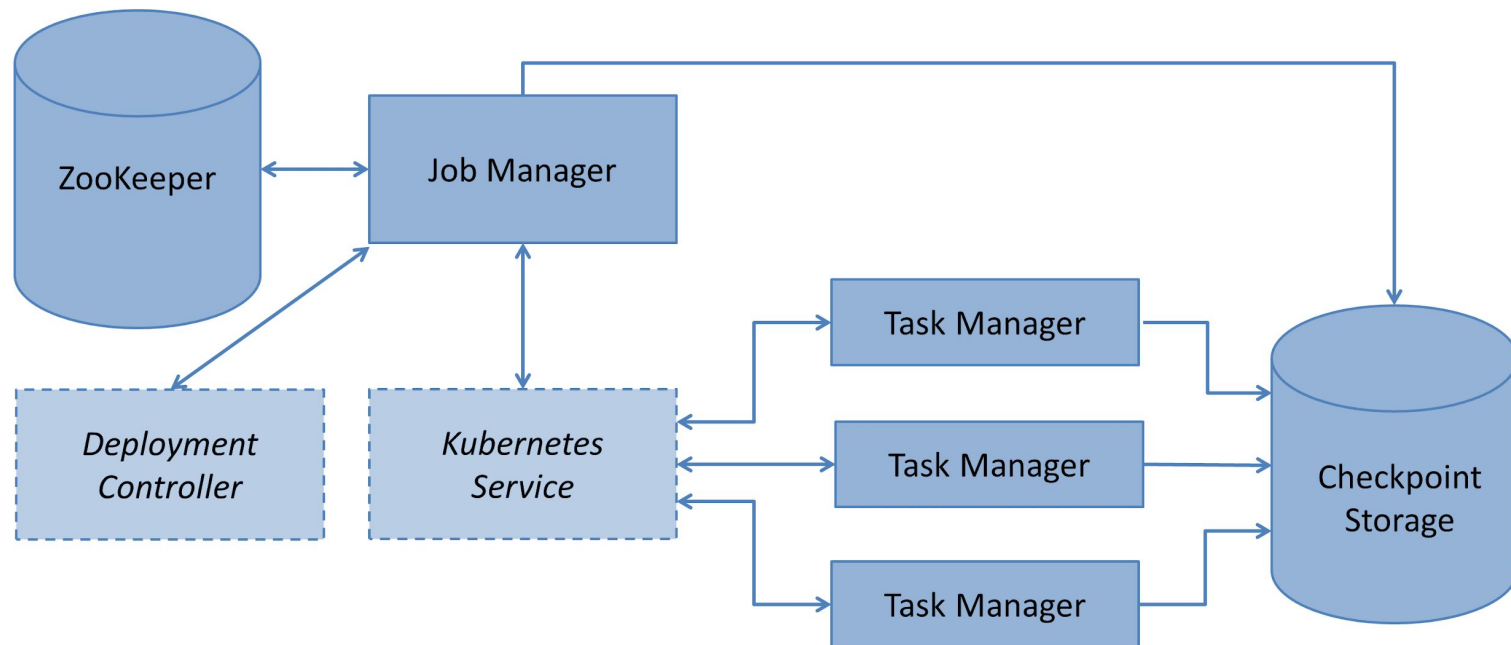
High Availability Mode (w/ Kubernetes or Yarn)



High Availability Mode (w/ Kubernetes or Yarn)



High Availability Mode (w/ Kubernetes or Yarn)



Our approach

- Making transformations stateful, enabling checkpointing in the applications and configuring restart strategies.
- Focus is on heterogeneous env, Multiple Job managers for high availability where one as leader and others as standby. Zookeeper can be used in case of Master(Job manager) failure to re-elect the Manager from Standby,
- To reduce latency we want the checkpoints/snapshots to be stored as close as possible to the node that fails,
- If possible add machine learning for adaptive fault tolerance.

Initial planning (next steps)

Cluster of Docker containers. First Single host for homogeneous env then multi host for heterogeneous env

Running flink cluster on those containers and run jobs.

Chaos testing using Pumba to check how application behaves in case of node or network failures.

Docker (incl. NETEM , cgroups (het. env))

Zookeeper, Kubernetes

Demo (Homogeneous Environment)

```
d071242@C02T427GHF1R ~> docker network ls
```

| NETWORK ID | NAME | DRIVER | SCOPE |
|--------------|--------|--------|-------|
| 488a762ee4e8 | bridge | bridge | local |
| c0e7560dc8d2 | flink | bridge | local |
| 267bb5c387c2 | host | host | local |

```
d071242@C02T427GHF1R ~/Flink> docker-compose up --scale taskmanager=3
```

```
flink_jobmanager_1 is up-to-date
```

```
Starting flink_taskmanager_1 ... done
```

```
Creating flink_taskmanager_2 ... done
```

```
Creating flink_taskmanager_3 ... done
```

```
Attaching to flink_jobmanager_1, flink_taskmanager_1, flink_taskmanager_3, flink_taskmanager_2
```

```
d071242@C02T427GHF1R ~> docker network ls
```

| NETWORK ID | NAME | DRIVER | SCOPE |
|--------------|---------------|--------|-------|
| 488a762ee4e8 | bridge | bridge | local |
| c0e7560dc8d2 | flink | bridge | local |
| dadd2c02367e | flink_default | bridge | local |
| 267bb5c387c2 | host | host | local |

```
d071242@C02T427GHF1R ~> docker exec -t -i flink_jobmanager_1 flink run /opt/flink/examples/streaming/SocketWindowWordCount.jar --hostname 10.183.176.71 --port 9000
```

Starting execution of program

Program execution finished

Job with JobID a5003bad6cf9efe11efe153de4727f2c has finished.

Job Runtime: 1371036 ms

```

"Containers": {
  "2e8caef8ca37dbff59011bc714549842432b991f9bdfbf9c22c50d823628dff8": {
    "Name": "flink_taskmanager_1",
    "EndpointID": "d8f3c8484c70dae0065d7c6a9d8357a7505134b33a2027062d96ba17dfceff18",
    "MacAddress": "02:42:ac:1f:00:03",
    "IPv4Address": "172.31.0.3/16",
    "IPv6Address": ""
  },
  "53b8cfecb5c8db25953d0d2b29ad818dae9559831f92fc5f352bddf8e92a3f2f": {
    "Name": "flink_taskmanager_2",
    "EndpointID": "036af5632cce8c3892357e64fcd933a6091399f98ae9531760c86f24429d8088",
    "MacAddress": "02:42:ac:1f:00:05",
    "IPv4Address": "172.31.0.5/16",
    "IPv6Address": ""
  },
  "84d14c4c2485b44584140773e9a8a121a3c93ce484b670432289ec0f892c648d": {
    "Name": "flink_jobmanager_1",
    "EndpointID": "cd36e9970c6dc9b66ab5b73437059dee42918030d671f20e7546c04b8090c115",
    "MacAddress": "02:42:ac:1f:00:02",
    "IPv4Address": "172.31.0.2/16",
    "IPv6Address": ""
  },
  "cc90bcfa728f956f2cf2f030baca45c6e0c46257e796200bc03db9c5cbae8d0f": {
    "Name": "flink_taskmanager_3",
    "EndpointID": "04fcbc5bdebca1259109de725f5c20a4cb610e2be7655dcaa5ab9b137c1755a2",
    "MacAddress": "02:42:ac:1f:00:04",
    "IPv4Address": "172.31.0.4/16",
    "IPv6Address": ""
  }
}

```



```
d071242@C02T427GHF1R ~> nc -l 9000
This is to test if wordcount is working properly
More texts to test
test test
d071242@C02T427GHF1R ~> nc -l 9000
Hello keep testing please
d071242@C02T427GHF1R ~> nc -l 9000
Some more text to test
We would like to check how it behaves when we stop the containers using Pumba
d071242@C02T427GHF1R ~> nc -l 9000
Now we will write more stuff to check how it executes the job when TM2 has been stopped
Keep writing
hello hello
```

```
d071242@C02T427GHF1R ~> pumba stop flink_taskmanager_1
d071242@C02T427GHF1R ~> pumba stop flink_taskmanager_2
```

```
taskmanager_1 | This : 1
taskmanager_1 | properly : 1
taskmanager_1 | working : 1
taskmanager_1 | wordcount : 1
taskmanager_1 | if : 1
taskmanager_1 | test : 1
taskmanager_1 | to : 1
taskmanager_1 | is : 2
taskmanager_1 | More : 1
taskmanager_1 | test : 3
taskmanager_1 | to : 1
taskmanager_1 | texts : 1
```

```
flink_taskmanager_1 exited with code 143
```

```
taskmanager_2 | Hello : 1
taskmanager_2 | please : 1
taskmanager_2 | testing : 1
taskmanager_2 | keep : 1
```

```
taskmanager_2 | Some : 1
taskmanager_2 | test : 1
taskmanager_2 | to : 1
taskmanager_2 | text : 1
taskmanager_2 | more : 1
taskmanager_2 | We : 1
taskmanager_2 | Pumba : 1
taskmanager_2 | using : 1
taskmanager_2 | containers : 1
taskmanager_2 | the : 1
taskmanager_2 | stop : 1
taskmanager_2 | we : 1
taskmanager_2 | when : 1
taskmanager_2 | behaves : 1
taskmanager_2 | it : 1
taskmanager_2 | how : 1
taskmanager_2 | check : 1
taskmanager_2 | to : 1
taskmanager_2 | like : 1
taskmanager_2 | would : 1
```

```
flink_taskmanager_2 exited with code 143
```

```
taskmanager_3 | Now : 1
taskmanager_3 | writing : 1
taskmanager_3 | Keep : 1
taskmanager_3 | stopped : 1
taskmanager_3 | been : 1
taskmanager_3 | has : 1
taskmanager_3 | TM2 : 1
taskmanager_3 | when : 1
taskmanager_3 | job : 1
taskmanager_3 | the : 1
taskmanager_3 | executes : 1
taskmanager_3 | it : 1
taskmanager_3 | how : 1
taskmanager_3 | check : 1
taskmanager_3 | to : 1
taskmanager_3 | stuff : 1
taskmanager_3 | more : 1
taskmanager_3 | write : 1
taskmanager_3 | will : 1
taskmanager_3 | we : 1
taskmanager_3 | hello : 2
```

```
"Containers": {  
  "84d14c4c2485b44584140773e9a8a121a3c93ce484b670432289ec0f892c648d": {  
    "Name": "flink_jobmanager_1",  
    "EndpointID": "cd36e9970c6dc9b66ab5b73437059dee42918030d671f20e7546c04b8090c115",  
    "MacAddress": "02:42:ac:1f:00:02",  
    "IPv4Address": "172.31.0.2/16",  
    "IPv6Address": ""  
  },  
  "cc90bcfa728f956f2cf2f030baca45c6e0c46257e796200bc03db9c5cbae8d0f": {  
    "Name": "flink_taskmanager_3",  
    "EndpointID": "04fcbc5bdebca1259109de725f5c20a4cb610e2be7655dcaa5ab9b137c1755a2",  
    "MacAddress": "02:42:ac:1f:00:04",  
    "IPv4Address": "172.31.0.4/16",  
    "IPv6Address": ""  
  }  
},
```

Interesting Points

When Pumba stops the container netcat listening on port 9000 exits automatically.

Job manager shows “Association failed” Caused by: [No route to host] which ultimately switches the job from RUNNING to FAILED.

When we re-run `nc -l 9000` and write something. We don't see anything in the logs. However, the moment we re-run the job Job manager assigns the task to another available task manager which spits this word count output.

Questions?